


Beating Atari with Natural Language guided Reinforcement Learning

Montezuma Revenge


Instructions to RL Agent

- 1: "Climb down the ladder"
- 2: "Jump to the rope"
- 3: "Go to the right side of the room"
- 4: "Climb down the ladder"
- 5: "Go to the left side of the room"
- 6: "Climb up the ladder"
- 7: "Get the key"




Implementation Details

Dataset Generation




Climb up the rope



Jump to the rope



Model Architecture




Bimodal Embeddings



● Climb up the ladder = 0.9527

Bimodal Embeddings



● Jump to the rope = 0.8666

Bimodal Embeddings




● Use the key = 0.8219

Bimodal Embeddings



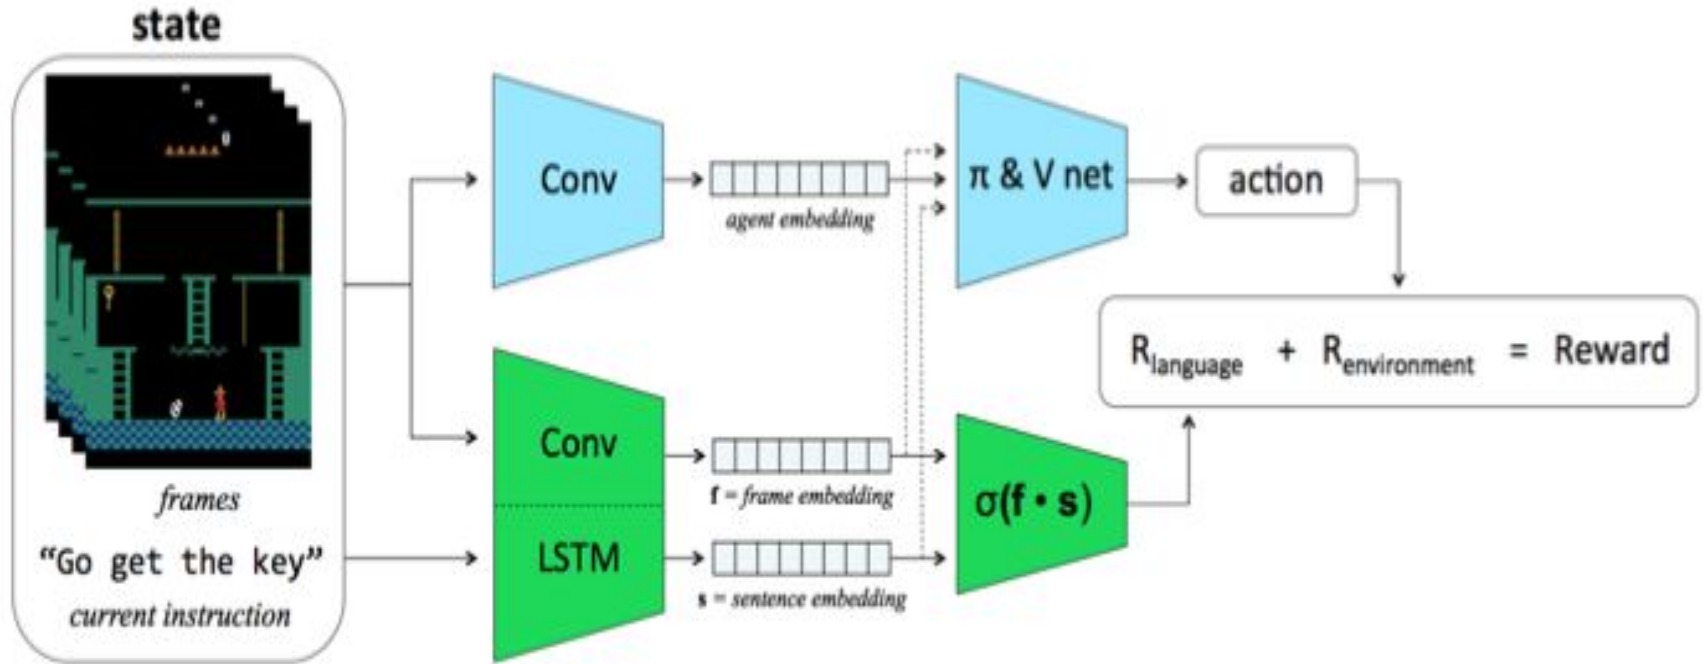
● Climb up the ladder = -0.2920

Bimodal Embeddings




- Go to the center of the room = 0.3710

Model Architecture




Simpler Games

MountainCar




CartPole



Problems Faced

- Action from pair of frames



Problems Faced

- Processing of sequence of instructions
 1. climb down the ladder
 2. jump to the rope
 3. go to the right side of the room
 4. climb down the ladder
 5. go to the bottom of the room
 6. go to the center of the room
 7. go to the left side of the room
 8. climb up the ladder
 9. get the key

